

## 6.1: MongoDB

**Instructor:** Dr. GP Saggese, [gsaggese@umd.edu](mailto:gsaggese@umd.edu)

- **References:**

- All concepts in slides
- MongoDB tutorial
- Web
  - <https://www.mongodb.com/>
  - Official docs
- Seven Databases in Seven Weeks, 2e

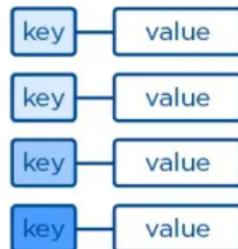


# Key-Value Store vs Document DBs

- **Key-value stores**

- Function as a map or dictionary
  - Examples: HBase, Redis
- Primarily retrieve values using keys
- Occasionally search within value fields using patterns
- Store uninterpreted values (e.g., binary blobs) linked to keys
- Use a single namespace for all key-value pairs

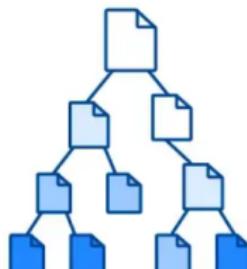
## Key-Value



- **Document DBs**

- Group key-value pairs into *documents*
  - Examples: MongoDB, CouchDB
- Documents formatted in JSON, XML, or BSON (binary JSON)
- Documents are part of *collections*
  - Comparable to *tables* in relational databases
- Large collections can be partitioned and indexed

## Document



# MongoDB

---



- Developed by MongoDB Inc
  - Founded in 2007
  - Based on DoubleClick's experience with large-scale data
  - Mongo comes from “hu-mongo-us”
- Highly popular NoSQL database
- Document-oriented NoSQL DB
  - Schema-less
    - No Data Definition Language (DDL) like SQL
    - Stores maps with any keys and values
    - Application manages schema, linking documents to meanings
  - Keys are string-stored hashes
    - Each document has a unique `_id` (reserved by MongoDB)
  - Values in BSON format
    - Based on JSON (B for Binary)
- High-performance
  - Developed in C++
  - Supports APIs (drivers) in various languages
    - Examples: JavaScript, Python, Ruby, Java, Scala, C++, etc

# MongoDB: Example of Document

---

- A document is a **JSON data structure**
- It corresponds to **a row in a relational DB**
  - Without schema
  - Primary key is `_id`
  - Values can be nested to an arbitrary depth

```
{  
    "_id" : ObjectId("4d0b6da3bb30773266f39fea"),  
    "country" : {  
        "$ref" : "countries",  
        "$id" : ObjectId("4d0e6074deb8995216a8309e")  
    },  
    "famous_for" : [  
        "beer",  
        "food"  
    ],  
    "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",  
    "mayor" : {  
        "name" : "Ted Wheeler",  
        "party" : "D"  
    },  
    "name" : "Portland",  
    "population" : 582000,  
    "state" : "OR"  
}
```

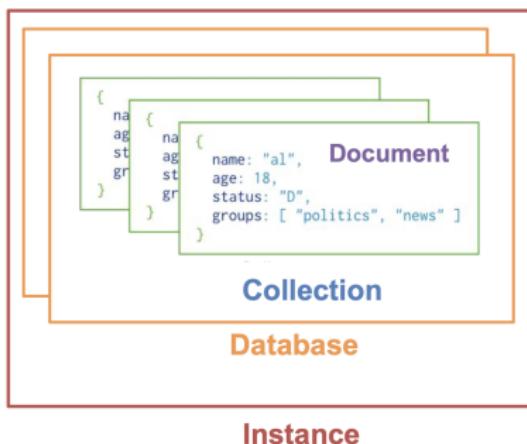
# MongoDB: Functionalities

---

- **Design goals**
  - Performance
  - Availability/scalability
  - Rich data storage (not rich querying!)
- **Dynamic schema**
  - No DDL
  - Secondary indexes
  - Query language via API
- **Several levels of data consistency**
  - Atomic writes and fully-consistent reads (document level)
- **No joins nor transactions across multiple documents**
  - Distributed queries easy and fast
- **High availability through replica sets**
  - Primary replication with automated failover
- **Built-in sharding**
  - Horizontal scaling via automated range-based partitioning
  - Reads and writes distributed over shards

# MongoDB: Hierarchical Objects

- A Mongo **instance** has:
  - Zero or more “databases”
  - Mongo instance ~ Postgres instance
- A Mongo **database** has:
  - Zero or more “collections”
    - Mongo collection ~ Postgres tables
  - Mongo database ~ Postgres database
- A Mongo **collection** has:
  - Zero or more “documents”
    - Mongo document ~ Postgres rows
- A Mongo **document** has:
  - One or more “fields”
    - Always has a primary key `_id`
    - Mongo field ~ Postgres columns



# Relational DBs vs MongoDB: Concepts

RDBMS Concept	MongoDB Concept	Meaning in MongoDB
database	database	Container for collections
relation / table / view	collection	Group of documents
row / instance	document	Group of fields
column / attribute	field	A name-value pair
index	index	Automatic
primary keys	_id field	Always the primary key
foreign key	reference	Pointers
table joins	embedded documents	Nested name-value pairs

```
{  
    "_id" : ObjectId("4d0b6da3bb30773266f39fea"),  
    "country" : {  
        "$ref" : "countries",  
        "$id" : ObjectId("4d0e6074deb8995216a8309e")  
    },  
    "famous_for" : [  
        "beer",  
        "food"  
    ],  
    "last_census" : "Sun Jan 07 2018 00:00:00 GMT -0700 (PDT)",  
    "mayor" : {  
        "name" : "Ted Wheeler",  
        "party" : "D"  
    },  
    "name" : "Portland",  
    "population" : 582000,  
    "state" : "OR"  
}
```

# Relational vs Document DB: Workflows

---

- **Relational DBs**

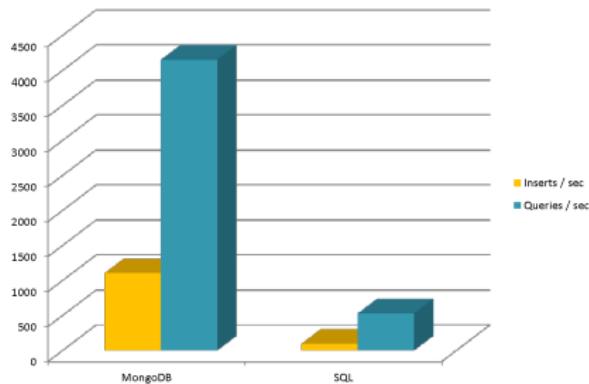
- E.g., PostgreSQL
- Know what to store
  - Tabular data
- Static schema allows query flexibility (e.g., joins)
- Complexity at insertion time
  - Decide data representation (schema)

- **Document DBs**

- E.g., MongoDB
- No assumptions on storage
  - E.g., irregular JSON data
- Access data by key
  - Nested key-value map
- Complexity at access time
  - Retrieve data from server
  - Process data client-side

# Why Use MongoDB?

- Simple and powerful to query
- Fast
  - 2-10x faster than PostgreSQL
- Data model suitable for most web applications
  - Semi-structured data
  - Quickly evolving systems
- Not suited for heavy, complex transaction systems
  - E.g., banking systems



# MongoDB: Data Model

- **Documents** are field-value pairs
  - *Field names*: strings
  - *Values*: any BSON type
    - Arrays of documents
    - Native data types
    - Other documents
- **Examples:**
  - `_id: ObjectId`
  - `name: document with fields first and last`
  - `birth and death: date type`
  - `contribs: array of strings`
  - `views: NumberLong type`

```
{  
    name: "sue",  
    age: 26,  
    status: "A",  
    groups: [ "news", "sports" ]  
}
```



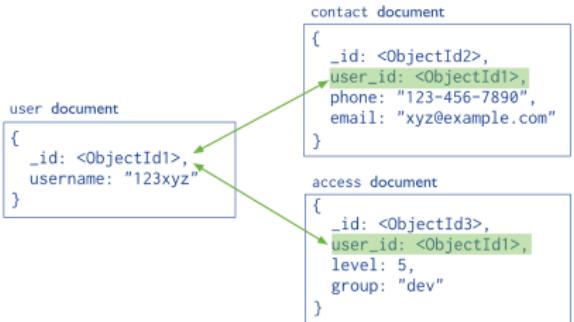
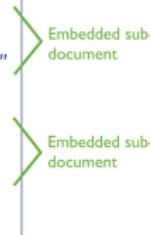
The diagram illustrates a MongoDB document structure. It consists of a left brace '{' followed by four key-value pairs. To the right of each pair is a double-headed horizontal arrow pointing between the field name and its value, indicating they are both part of the same document object.

```
{  
    _id: ObjectId("5099883df3f4948bd2f98391"),  
    name: { first: "Alan", last: "Turing" },  
    birth: new Date('Jun 23, 1912'),  
    death: new Date('Jun 07, 1954'),  
    contribs: [ "Turing machine", "Turing test", "Turingery" ],  
    views : NumberLong(1250000)  
}
```

# MongoDB: Data Model

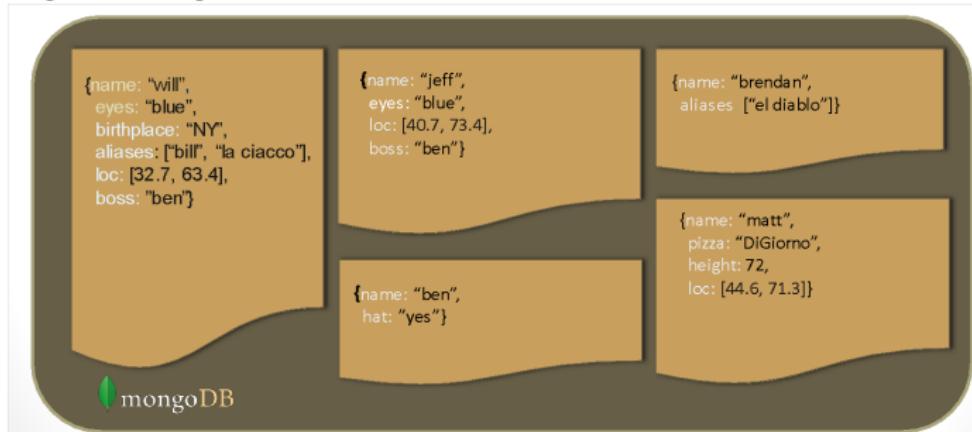
- **Documents can be nested**
  - Embedded sub-document
- **Denormalized data models**
  - Store related information in the same record
  - Avoids the need for a join operation
- **Normalized data models**
  - Eliminate duplication
  - Represent many-to-many relationships

```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



# Schema Free

- MongoDB does not need pre-defined data schema
- Every **document** in a **collection** can have different fields and values
  - No need for NULL values / union of fields like in relational DBs
- E.g., heterogeneous data instances



# JSON Format

---

- **JSON = JavaScript Object Notation**

- Data is stored in field/value pairs
- A field/value pair consists of:
  - A field name (always a string)
  - Followed by a colon :
  - Followed by a typed value

`"name": "R2-D2"`

- Just references
- Data in documents is separated by commas ,

`"name": "R2-D2", "race": "Droid"`

- Curly braces {} hold documents

`{ "name": "R2-D2", "race": "Droid", "affiliation": "rebels" }`

- An array is stored in brackets []

`[`

`{ "name": "R2-D2", "race": "Droid", "affiliation": "rebels" },  
  { "name": "Yoda", "affiliation": "rebels" }`

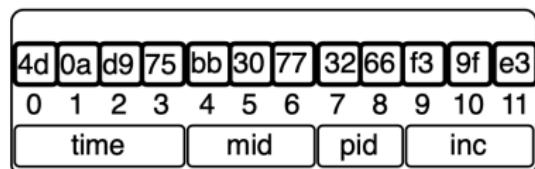
# BSON Format

---

- **Binary-encoded serialization of JSON-like documents**
  - <https://bsonspec.org>
  - Similar to Protocol Buffers, but more schema-less
- **Optimized for random access**
  - Prefixes elements in a BSON document with a length field
- **MongoDB understands BSON objects, even nested ones**
  - Builds indexes and matches objects against query expressions for BSON keys

# ObjectId

- Each JSON document contains an `_id` field of type ObjectId
  - Similar to the SERIAL constraint incrementing a numeric primary key in PostgreSQL
- An ObjectId is 12 bytes, composed of:
  - Timestamp
  - Client machine ID
  - Client process ID
  - 3-byte auto-incremented counter
- Each MongoDB process handles its own ID generation without collisions
  - Due to MongoDB's distributed nature
- Details [here](#)



# Indexes

---

- Primary index
  - Automatically created on the `_id` field
  - B+ tree indexes
- Secondary index
  - Improve query performance
  - Can enforce unique values for a field
- Single field index and compound index (like SQL)
  - Order of fields in a compound index matters
- Sparse property of an index
  - Index contains entries only for documents with the indexed field
  - Ignores records without the field
- Rejects records with duplicate keys if the index is unique and sparse
- Details [here](#)