



## UMD DATA605 - Big Data Systems

### 7.3: Serialization Formats

- **Instructor:** Dr. GP Saggese - [gsaggese@umd.edu](mailto:gsaggese@umd.edu)

# Serialization Formats

---

- Programs need to send data to each other (network, disk)
  - E.g., remote Procedure Calls (RPCs)
  - Recent technologies based on schemas
    - JSON, YAML, Protocol Buffer, Python Pickle
- Serialization formats are data models

# Comma Separated Values (CSV)

- CSV stores data row-wise as text without schema
  - Each line is a data record
  - Records have fields separated by commas
- **Pros**
  - Very portable
    - Text format
    - Supported by all tools
  - Human-friendly

- **Cons**
  - Large footprint
    - Requires compression
  - Parsing is CPU intensive
  - No easy random access
  - Can't read only a subset of columns
  - No schema/types
    - Annotate CSV with schema
  - Mainly read-only, hard to modify

| Year | Make  | Model                                  | Description                          | Price   |
|------|-------|--|--------------------------------------|---------|
| 1997 | Ford  | E350                                   | ac, abs, moon                        | 3000.00 |
| 1999 | Chevy | Venture "Extended Edition"             |                                      | 4900.00 |
| 1999 | Chevy | Venture "Extended Edition, Very Large" |                                      | 5000.00 |
| 1996 | Jeep  | Grand Cherokee                         | MUST SELL!<br>air, moon roof, loaded | 4799.00 |

```
Year,Make,Model,Description,Price
1997,Ford,E350,"ac, abs, moon",3000.00
1999,Chevy,"Venture ""Extended Edition""","",4900.00
1999,Chevy,"Venture ""Extended Edition, Very Large""","",5000.00
1996,Jeep,Grand Cherokee,"MUST SELL!
air, moon roof, loaded",4799.00
```

# (Apache) Parquet

---



- Parquet reads data tiles
- Supports multi-dimensional, nested data
  - Generalizes dataframes
- Column-storage
  - Stores each column together, uniform data type, compressed efficiently
- IO layer executes queries
  - Reads only necessary data chunks from disk
- **Pros**
  - 10x smaller than CSV
  - 10x faster with multi-threading
  - Read subset of columns and rows
- **Cons**
  - Binary, not human-friendly
  - Requires ingestion step to convert to Parquet
  - Mainly read-only, hard to modify

# JSON

---

- JSON = JavaScript Object Notation
- Nested dictionaries and arrays
- Similar to XML
  - More human-readable
  - Less boilerplate
  - Executable in JavaScript and Python

```
{  
    "firstName": "John",  
    "lastName": "Smith",  
    "isAlive": true,  
    "age": 25,  
    "height_cm": 167.6,  
    "address": {  
        "streetAddress": "21 2nd Street",  
        "city": "New York",  
        "state": "NY",  
        "postalCode": "10021-3100"  
    },  
    "phoneNumbers": [  
        {  
            "type": "home",  
            "number": "212 555-1234"  
        },  
        {  
            "type": "office",  
            "number": "646 555-4567"  
        }  
    ],  
    "children": [],  
    "spouse": null  
}
```

# Protocol Buffers

---

- Developed by Google
- Open-source
- Represent data structures:
  - Language agnostic
  - Platform agnostic
  - Versioning
- Schema is mostly relational
  - Optional fields
  - Types
  - Default values
  - Structures
  - Arrays
- Schema specified using a .proto file
- Compiled by protoc to produce C++, Java, or Python code to initialize, read, serialize objects

```
import addressbook_pb2
person = addressbook_pb2.Person()
person.id = 1234
person.name = "John Doe"
person.email = "jdoe@example.com"
phone = person.phones.add()
phone.number = "555-4321"
phone.type = addressbook_pb2.Person.HOME
message Person {
    optional string name = 1;
    optional int32 id = 2;
    optional string email = 3;
    enum PhoneType {
        MOBILE = 0;
        HOME = 1;
        WORK = 2;
    }
    message PhoneNumber {
        optional string number = 1;
        optional PhoneType type = 2;
    }
    repeated PhoneNumber phones = 4;
}
```



# Serialization Formats

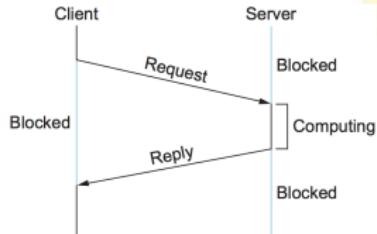
---

- Avro
  - Richer data structures
  - JSON-specified schema
- Thrift
  - Developed by Facebook
  - Now Apache project
  - More languages supported
  - Supports exceptions and sets

```
{  
  "namespace": "example.avro",  
  "type": "record",  
  "name": "User",  
  "fields": [  
    {  
      "name": "name",  
      "type": "string"  
    },  
    {  
      "name": "favorite_number",  
      "type": [  
        "int",  
        "null"  
      ]  
    },  
    {  
      "name": "favorite_color",  
      "type": [  
        "string",  
        "null"  
      ]  
    }  
  ]  
}
```

# Remote Procedure Call

- **Remote Procedure Call (RPC)** requests services from programs on other computers, abstracting network communication
- **Goal:** Make remote calls like local procedure calls without network details
- **Problems**
  - Can't serialize pointers
  - Asynchronous communication
  - Failures and retry
- Used in distributed systems
  - E.g., microservices, cloud services, client-server applications
- Can be synchronous or asynchronous



# RPCs: Internals

- *Client procedure call:* Client calls stub function with arguments
- *Request marshalling:* Client stub serializes arguments for network transmission
- *Server communication:* Client's RPC runtime sends request to server
- *Server-side unmarshalling:* Server's RPC runtime deserializes arguments
- *Procedure execution:* Server calls procedure
- *Response marshalling:* Return values marshaled into response message
- *Client communication / response unmarshalling / return to client:* Return values passed back to client's stub call, execution

