





## UMD DATA605 - Big Data Systems

### 8.1: Cluster Architecture

- **Instructor:** Dr. GP Saggese - [gsaggese@umd.edu](mailto:gsaggese@umd.edu)
- **Resources**
  - Silberschatz: Chap 10



# Big Data: Storing and Computing




- Big data **needs 10k-100k machines**
- **Two problems**
  - Storing big data
  - Processing big data
- Need to **solve together and efficiently**
  - One slow phase slows entire system

# Processing the Web: Example

---

- Web contains (in 2024):
  - 20+ billion pages
  - 5M TBs
- Need 1M 5TB hard drives
  - \$100/HDD -> \$100M in total
- One computer reads 300 MB/sec
  - 4,433 years to read web serially!
- Much more time needed for data processing!



# How to Store Big Data?

---

- There are many different solutions to storing big data?
  1. **Distributed file systems**
  2. **Sharding across multiple DBs**
  3. **Parallel and distributed DBs**
  4. **Key-value stores**

# 1) Distributed File Systems

---

- Files stored across machines, single file-system view to clients
  - E.g., Google File System (GFS)
  - E.g., Hadoop File System (HDFS) based on GFS
  - E.g., AWS S3
- Files are:
  - Broken into blocks
  - Blocks partitioned across machines
  - Blocks often replicated
- **Goals:**
  - Store data not fitting on one machine
  - Increase performance
  - Increase reliability/availability/fault tolerance

## 2) Sharding Across Multiple DBs

---

- **Sharding:** Partition records across multiple DBs or machines
  - Shard keys, aka partitioning keys / partition attributes
  - Range partition (e.g., timeseries)
  - Hash partition
- **Pros**
  - Scale beyond centralized DB for more users, storage, processing speed
- **Cons**
  - Replication needed for failures
  - Ensuring consistency is challenging
  - Relational DBs struggle with constraints (e.g., foreign key) and transactions on multiple machines

### **3) Parallel and Distributed DBs**

---

- Store and process data on multiple machines (cluster)

- **Pros**

- From programmer viewpoint
  - Traditional relational DB interface
  - Appears as a single-machine DB
- Scale to 10s-100s of machines
- Data replication enhances performance and reliability
  - Frequent failures with 100s of machines
  - Queries can restart on different machines

- **Cons**

- Incremental query execution is complex
- Scalability limits

## 4) Key-value Stores

---

- **Problem**

- Applications store billions of small records
- RDBMSs lack multi-machine constraints and transactions

- **Solution**

- Key-value stores / Document / NoSQL systems
- Store, update, retrieve records by key

- **Pros**

- Partition data across machines
- Support replication and consistency
- Balance workload, add machines

- **Cons**

- Sacrifice features for scalability
  - Declarative querying
  - Transactions
  - Non-key attribute retrieval

# 4 Parallel Key-value Stores

---

- **Parallel key-value stores**
  - Redis
  - Google BigTable
  - Apache HBase (open source BigTable)
  - AWS Dynamo, S3
  - Cassandra (Facebook)
  - Azure cloud storage (Microsoft)
- **Parallel document stores**
  - MongoDB cluster
  - Couchbase
- **In-memory caching systems**
  - Store relations in-memory
  - Replicated or partitioned across machines
  - E.g., memcached or Redis

# How to Compute with Big Data?

---

- **Challenges**


- Distribute computation
- Simplify writing distributed programs
  - Distributed/parallel programming is hard
- Store data in a distributed system
- Survive failures
  - One server lasts ~3 years (1,000 days)
  - With 1,000 servers, expect 1 failure/day
  - E.g., 1M machines (Google in 2011) → 1,000 machines fail daily

- **MapReduce**

- Solve problems for specific computations
- Elegant way to work with big data
- Originated as Google's data manipulation model
  - Not an entirely new idea

# Cluster Architecture

- Standard architecture for big data computation:
  - Cluster of commodity Linux nodes
  - Commodity network (typically Ethernet) to connect nodes
  - 2011: Google ~1M machines
  - 2025: ~10-15M (?)



# Cluster Architecture

---



# Cluster Architecture: Network Bandwidth

---

- **Problems**

- Data hosted on different machines
- Network data transfer takes time

- **Solutions**

- Bring computation to data
- Store files multiple times for reliability/performance

- **MapReduce**

- Addresses these problems
- Storage: distributed file system
  - Google GFS, Hadoop HDFS
- Programming model: MapReduce

# Storage Infrastructure


---

- **Problem**
  - Store data persistently and efficiently despite node failures
- **Typical data usage pattern**
  - Huge files (100s of GB to 1TB)
  - Common operations: reads and appends
  - Rare in-place updates
- **Solution**
  - Distributed file system
  - Store files across multiple machines
  - Files are:
    - Broken into blocks
    - Partitioned across machines
    - Replicated across machines
  - Provide a single file-system view to clients

# Distributed File System

- Reliable distributed file system
  - Data in “chunks” across machines
  - Each chunk replicated on different machines
  - Seamless recovery from disk or machine failure
- Bring computation directly to the data
  - “chunk servers” also serve as “compute servers”

2-10 Gbps backbone  
between racks



# Hadoop Distributed File System

---

- **NameNode**

- Store file/dir hierarchy
- Store file metadata (location, size, permissions)

- **DataNodes**

- Store data blocks
- Split file into 16-64MB blocks
- Replicate chunks (2x or 3x)
- Keep replicas in different racks

# Hadoop Distributed File System

---

- **Library for file access**
  - Read:
    - Contact *NameNode* for *DataNode* and block pointer
    - Connect to *DataNode* for data access
  - Write:
    - *NameNode* creates blocks
    - Assign blocks to multiple *DataNodes*
    - Client sends data to *DataNodes*
    - *DataNodes* store data
- **Client**
  - API (e.g., Python, Java) to library
  - Mount HDFS on local filesystem