



9.1: Apache Spark: Principles

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu
- **References:**
 - Concepts in the slides
 - Academic paper
 - “Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing”, 2012
 - Mastery
 - “Learning Spark: Lightning-Fast Data Analytics” (2nd Edition)
 - Not my favorite, but free here



Hadoop MapReduce: Shortcomings

- **Hadoop is hard to administer**
 - Many layers (HDFS, Yarn, Hadoop, ...)
 - Extensive configuration
- **Hadoop is hard to use**
 - Verbose API
 - Limited language support (e.g., Java is native)
 - MapReduce jobs read / write data on disk



- **Large but fragmented ecosystem**
 - No native support for:
 - Machine learning
 - SQL, streaming
 - Interactive computing
 - New systems developed on Hadoop for new workloads
 - E.g., Apache Hive, Storm, Impala, Giraph, Drill

(Apache) Spark

- Open-source
 - DataBrick monetizes it (\$40B startup)



- General processing engine

- Large set of operations beyond Map() and Reduce()
- Combine operations in any order
- Transformations vs Actions
- Computation organized as a DAG
 - DAGs decomposed into parallel tasks
- Scheduler/optimizer for parallel workers

- Supports several languages

- Java, Scala (preferred), Python supported through bindings

- Data abstraction

- Resilient Distributed Dataset (RDD)
- DataFrames, Datasets built on RDDs

- Fault tolerance through RDD lineage

In-Memory computation

Berkeley: From Research to Companies


- Amplab
- Rise lab



Berkeley AMPLab Data Analytics Stack

- So many tools that they have their own big data stack!

<https://amplab.cs.berkeley.edu/software/>




Apache Spark

- **Unified stack**
 - Different computation models in a single framework
- **Spark SQL**
 - ANSI SQL compliant
 - Work with structured relational data
- **Spark MLlib**
 - Build ML pipelines
 - Support popular ML algorithms
 - Built on top of Spark DataFrame
- **Spark Streaming**
 - Handle continually growing tables
 - Tables are treated as static table
- **GraphX**
 - Manipulate graphs
 - Perform graph-parallel computation
- **Extensibility**
 - Read from many sources
 - Write to many backends



One computation engine




General purpose applications



Resilient Distributed Dataset (RDD)

- **A Resilient Distributed Dataset (RDD)**

- Collection of data elements
- Partitioned across nodes
- Operated on in parallel
- Fault-tolerant
- In-memory / serializable



- **Applications**

- Best for applications applying the same operation to all dataset elements (vectorized)
- Less suitable for asynchronous fine-grained updates to shared state
 - E.g., updating one value in a dataframe

- **Ways to create RDDs**

- *Reference* data in external storage
 - E.g., file-system, HDFS, HBase
- *Parallelize* an existing collection in your driver program
- *Transform* RDDs into other RDDs


Transformations vs Actions

- **Transformations**

- Lazy evaluation
- Nothing computed until an Action requires it
- Build a graph of transformations

- **Actions**

- When applied to RDDs force calculations and return values
- Aka Materialize



Spark Example: Estimate Pi

```
# Estimate π (compute-intensive task).
# Pick random points in the unit square [(0,0)-(1,1)].
# See how many fall in the unit circle center=(0, 0), radius=1.
# The fraction should be π / 4.

import random
random.seed(314)

def sample(p):
    x, y = random.random(), random.random()
    in_unit_circle = 1 if x*x + y*y < 1 else 0
    return in_unit_circle

# "parallelize" method creates an RDD.
NUM_SAMPLES = int(1e6)
count = sc.parallelize(range(0, NUM_SAMPLES)) \
        .map(sample) \
        .reduce(lambda a, b: a + b)
approx_pi = 4.0 * count / NUM_SAMPLES
print("pi is roughly %f" % approx_pi)
```

executed in 386ms, finished 04:27:53 2022-11-23

pi is roughly 3.141400

SCIENCE
ACADEMY



Spark: Architecture

- **Architecture**

- Who does what
- Responsibilities of each piece

- **Spark Application**

- Code describing computation
- E.g., Python code calling Spark

- **Spark Driver**

- Instantiate *SparkSession*
- Communicate with *Cluster Manager* for resources
- Transform operations into DAG computations
- Distribute task execution across *Executors*

- **Spark Session**


- Interface to Spark system

- **Cluster Manager**

- Manage and allocate resources
- Support Hadoop, YARN, Mesos, Kubernetes


- **Spark Executor**

- Run worker node to execute tasks
- Typically one executor per node



Spark: Computation Model

- **Architecture** = who does what
- **Computational model** = how are things done




- **Spark Driver**
 - Converts Spark application into one or more Spark *Jobs*
 - Describes computation with *Transformations* and triggers with *Actions*
- **Spark Job**
 - Parallel computation in response to a Spark *Action*
 - Each *Job* is a DAG with one or more dependent *Stages*
- **Spark Stage**
 - Smaller operation within a *Job*
 - Stages can run serially or in parallel
- **Spark Task**
 - Each Stage has multiple *Tasks*

Deployment Modes

- Spark can run on several different configurations
 - **Local**
 - E.g., run on your laptop
 - Driver, Cluster Manager, Executors all run in a single JVM on the same node
 - **Standalone**
 - Driver, Cluster Manager, Executors run in different JVMs on different nodes
 - **YARN or Kubernetes**
 - Driver, Cluster Manager, Executors run on different pods (i.e., containers)


Distributed Data and Partitions

- **Data is distributed** as partitions across different physical nodes
 - Store each partition in memory
 - Enable efficient parallelism
- **Spark Executors** process data “close” to them
 - Minimize network bandwidth
 - Ensure data locality
 - Similar approach to Hadoop



Parallelized Collections

- Parallelized collections created by calling *SparkContext parallelize()* on an existing collection



- Data spread across nodes
- Number of *partitions* to cut dataset into
 - Spark runs one *Task* per partition
 - Aim for 2-4 partitions per CPU
 - Spark sets partitions automatically based on your cluster
 - Set manually by passing as a second parameter to `parallelize()`