9.3: Python Dask

- **Instructor**: Dr. GP Saggese, gsaggese@umd.edu
- **Resources**
  - Web resources:
    - Dask project
    - Dask examples
  - Tutorial
    - Dask_tutorial
    - Dask_advanced_tutorial
  - Class project
  - Mastery
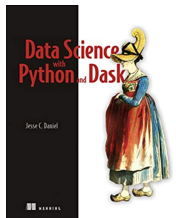   - Data science with Python and Dask, 2019

Figure 1: alt_text

SCIENCE
ACADEMY

Dataset Size Issues

- **Small datasets (< 1 GB)**
  - Fits into RAM
  - Manipulation doesn't require paging to disk
- **Medium dataset (< 1TB)**
  - Doesn't fit into RAM
  - Fits into local disk
    - Performance penalty imposed by using local disk
  - Need multiple CPU cores
    - Difficult to take advantage of parallelism with Python / Pandas
- **Large dataset (> 1TB)**
  - Doesn't fit into RAM
  - Doesn't fit into local disk
  - Need multiple

SCIENCE ACADEMY

Dataset Size Issues

- **Small datasets**
  - < 1 GB
- **Medium dataset**
  - < 1TB
- **Large dataset**
  - *1TB*
- **The thresholds are fuzzy and changing over time**
  - E.g., you can scale the computer 10x and get 10x bigger data sets
- **Problem with scaling datasets**
  - Long run times
  - Rewriting code in different language / API for datasets of different size
  - Need to think about *what to do* it and *how to do* it efficiently
  - Cumbersome framework (Pandas easy, Hadoop difficult)

SCIENCE
ACADEMY

Dask

- **Dask is written in Python**
  - It scales natively Numpy, Pandas, sklearn
  - Dask objects are wrappers (don't just mirror the interface) objects from the respective libraries (e.g., Pandas DataFrame, numpy array)
  - Parallel parts are called "chunks" or "partitions"
    - Are queued to be worked on
    - Shipped between machines
    - Worked locally on a machine

- **Pros**
  - Users don't need to learn a new language, but can use familiar interfaces
  - Can focus on writing code that is optimized for parallelism
    - Dask does the heavy lifting

- **Scaling Dask is easy**
  - Users can write a prototype task on a local machines and use a cluster when needed
  - No need to refactor existing code
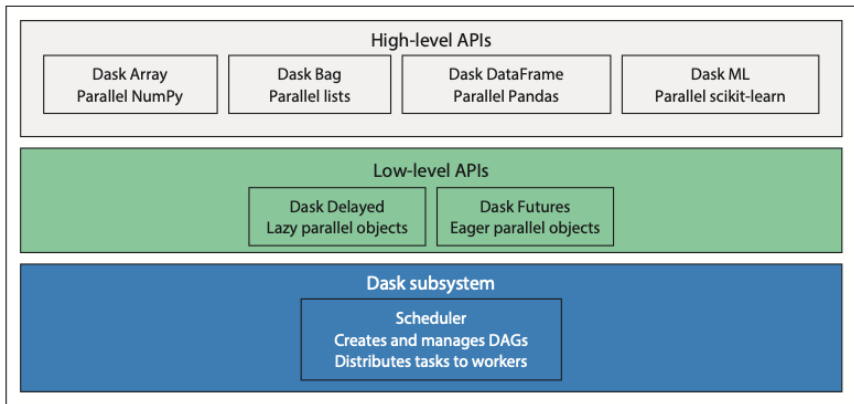

Figure 2: alt_text

Dask Layers

Figure 3: alt_text

Scaling Up vs Scaling Out

- **Scaling up**
  - = replace equipment with larger, faster equipments
    - E.g., buy a larger pot, replace knife with food processor
  - **Pros**
    - You got better hardware, nothing else needs to change (e.g., code)
  - **Cons**
    - There will be a time where you exceed the capacity of the current machines
    - Cost: more powerful machines are expensive
- **Scaling out**
  - = divide the work between many workers in parallel
    - E.g., buy more pots and hire more cooks
  - **Pros**
    - Task scheduler organizes computation, assigning workers to each task
    - More cost-effective solution since no specialized hardware is needed
  - **Cons**
    - Need to write code to expose parallelism
    - Costs of maintaining a cluster

Dask: Computation

- **Lazy computations**
  - User defines the transformations on the data
  - No need to wait for one computation to finish before defining the next
  - Avoid loading the entire data in memory by operating in chunks
  - E.g.,
    - Split a 2GB file into 32 64MB chunks
    - Operate on 8 chunks at a time on each server
    - The max memory consumption doesn't exceed 512MB $= (8 \times 32)$
  - Each task tracks object dimensions and data types
   - No code is executed
- **compute()**
  - Running a computation (aka materializing) "' python missing_count_pct = missing_count.compute()"'
- **persist()**
  - As soon as a node in the graph emits results, its intermediate work is
  - discarded to minimize memory usage
  - If we need to do additional computation on intermediate nodes we need to re-run the graph
  - **persist()** tells Dask to keep the intermediate result in memory
  - This speeds up a large and complex DAG that needs to be reused many

Dask: Data Structures

- **Dask DataFrame** implements Pandas DataFrame
  - Tabular / relational data
- **Dask Array** implements numpy ndarray
  - Multidimensional array**
- **Dag Bag** coordinates Python lists of objects
  - Parallelize computations on unstructured or semi-structured data
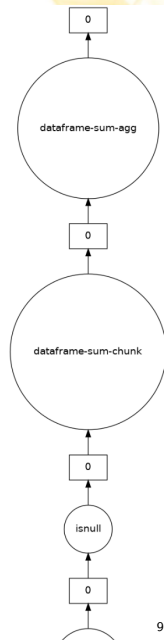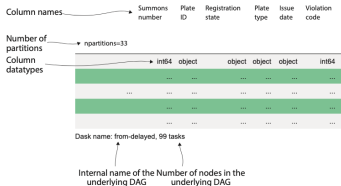

Pandas DataFrame

Dask DataFrame

NumPy Array

Dask Array

```
[1, 2, 3, 4, 5]
[1, 2, 3] [4, 5]
```

SCIENCE ACADEMY
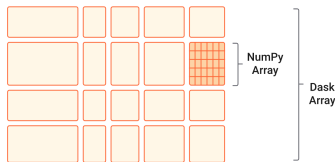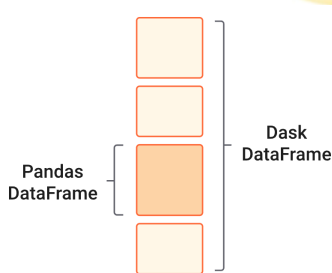
Dask Reading Data

- dask.dataframe.read_csv()
  - Doesn't load the data in memory with
  - Tries to infer the types of the columns
    - By randomly sampling some data
    - Best to set the data types
    - Even better is to use Parquet since it stores data and types together
- Partitions = chunks of data that can be worked independently
  - E.g., 33 partitions
  - Graph is composed of 99 tasks
  - Each partition reads data, splits data, initializes df object

Column names → Summons number, Plate ID, Registration state, Plate type, Issue date, Violation code

Number of partitions → npartitions=33

Column datatypes → int64, object, object, object, object, int64

Dask name: from-delayed, 99 tasks

Internal name of the underlying DAG — Number of nodes in the underlying DAG

dataframe-sum-agg

dataframe-sum-chunk

isnull

SCIENCE ACADEMY

Low Level APIs: Delayed

- Handle computations that don't fit in native Dask data structures (e.g., Dask DataFrame)
- In the example below there is parallelism that can be exploited

```python
def inc(x):
    return x + 1

def double(x):
    return x * 2

def add(x, y):
    return x + y

data = [1, 2, 3, 4, 5]

output = []
for x in data:
    # (x + 1) + (x * 2) = 3x + 1
    a = inc(x)
    b = double(x)
    c = add(a, b)
    # 1 -> 4
    # 2 -> 7
    # 3 -> 10
    # 4 -> 13
    # 5 -> 16
    output.append(c)

# 4 + 7 + 10 + 13 + 16 = 20 + 20 + 10 = 50
total = sum(output)
print(total)
```


SCIENCE
ACADEMY
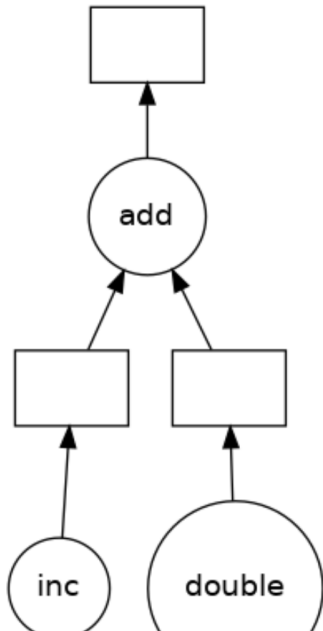
Low Level APIs: Futures

- In parallel programming, a "future" encapsulates the asynchronous execution of a callable, representing the eventual result of the operation
- Futures is the most general way of specifying concurrency in Dask
  - Everything can be expressed in terms of futures
  - User can specify what's blocking and what's not blocking
- Python**concurrent.futures**
  - High-level interface for asynchronously executing callables
  - Thread pool or Process
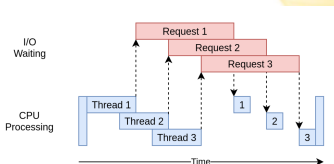

```python
def inc(x):
    return x + 1


def add(x, y):
    return x + y


a = client.submit(inc, 10)
b = client.submit(inc, 20)
```

```
>>> a
<Future: status: pending, key: inc-b8aaf26b99466a7a1980efa1ade6701d>
>>> a
<Future: status: finished, type: int, key: inc-b8aaf26b99466a7a1980efa1ade6701d>
>>> a.result()   # blocks until task completes and data arrives
11
```
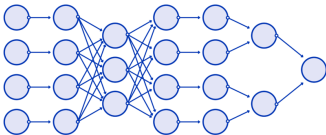
Different Types of Parallel Workload

- Break program in medium-size tasks of computation
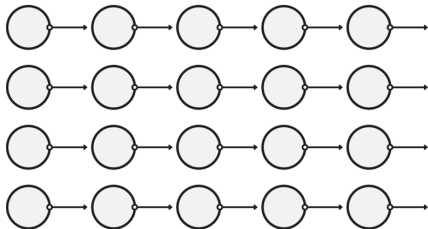 - E.g., a function call
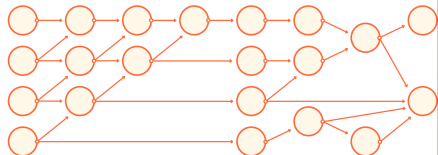
MapReduce
Hadoop/Spark/Dask

Embarrassingly Parallel
Hadoop/Spark/Dask/Airflow/Prefect

Full Task Scheduling
Dask/Airflow/Prefect

SCIENCE
ACADEMY

Encoding Task Graph

- Dask encodes tasks in terms of Python dicts and functions ::: columns ::::

```python
def inc(i):
    return i + 1


def add(a, b):
    return a + b


x = 1
y = inc(x)
z = add(y, 10)
```

SCIENCE
ACADEMY

Task Scheduling

- Data collections (Bags, Arrays, DataFrame) and their operations create task graphs
  - Nodes in the task graph are Python functions
  - Edges are dependencies (e.g., output from one task used as input in another task)
- Task graphs are scheduled for execution
- Single-machine scheduler
  - Use local process or thread pool
  - Simple but it can only run on a single machine
- Distributed scheduler
  - It can run locally or distributed across a cluster
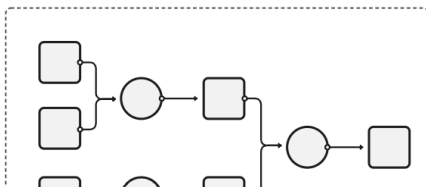
Collections
(create task graphs)

→

Task Graph

→

Schedulers
(execute task graphs)

Dask Array

Dask DataFrame

Dask Bag

Single-machine (threads, processes, synchronous)

Distributed

Task Scheduling

- **Dask task scheduler orchestrates the work dynamically**
  - Not a static scheduling of operations like a relational DB
  - When the computation takes place, Dask dynamically assesses:
    - What tasks has been completed
    - What tasks is left to do
    - What resources (CPUs) are free
    - Where the data is located
- **This dynamic approach handles a variety issues:**
  - Worker failure
    - Just re-run
  - Workers completing work at different speeds

SCIENCE
ACADEMY
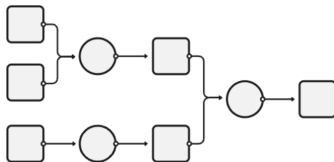
Dask vs Spark

- Spark has
- **Pros**
  - Popular framework for analyzing large datasets
  - In-memory alternative to MapReduce / Hadoop
- **Cons**
  - Spark is a Java library, supporting Python through PySpark API
    - Python code is executed on JVM through py4j
    - Difficult to debug since execution occurs outside Python
  - Different DataFrame API than Pandas
    - Learn how to do things "the Spark way"
    - You might need to implement things twice to go from exploratory analysis to large experiments / production
  - Optimized for MapReduce operations over a collection
  - Difficult to set-up and configure

SCIENCE
ACADEMY

Tutorial

Tutorial - From the official documentation
https://docs.dask.org/en/stable/10-minutes-to-dask.html