



## UMD DATA605 - Big Data Systems

### 12.1: Streaming and Real-time Analytics

- **Instructor:** Dr. GP Saggese, [gsaggese@umd.edu](mailto:gsaggese@umd.edu)

# Motivation

---

- Often Big Data is generated continuously
- Applications grow daily, e.g.,
  - Financial data
  - Sensors, RFID, IoT
  - Network/systems monitoring
  - Video/audio data
- Data stream is a sequence of events in time
- Stream processing is a paradigm for data streams as computation objects
- Needs to support
  - High data rates
  - Real-time processing with low latencies
    - Time-series operations
  - Data dissemination
    - Send data subset to interested users
  - Distributed operation
    - Single node often insufficient for real-time processing

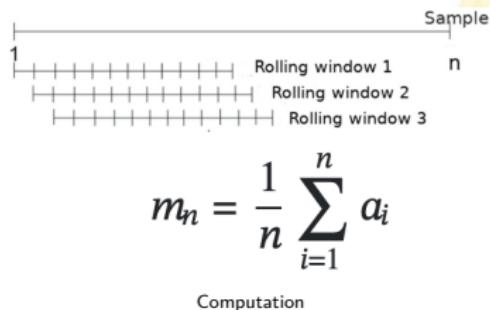
# Examples of Data Stream Tasks

---

- Any SQL query can be continuous
  - E.g., "compute moving average over last hour every 10 mins"
  - StreamSQL extends SQL to support windows over streams
- Pattern recognition
  - E.g., "alert me when A occurs and then B within 10 mins"
- Surveillance
- Fraud detection and prevention
  - Anti-money laundering
  - RegNMS, MiFID
- Financial modeling
- Statistical tasks
  - E.g., de-noising measured readings
  - Build an online machine learning model
- Process multimedia data
  - E.g., online object detection, activity detection

# Why Not Using the Usual RDBM?

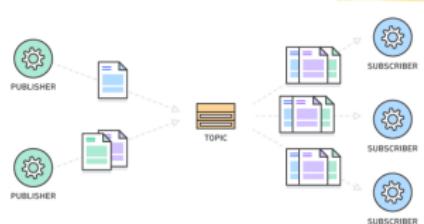
- Simple case: “*report moving average over last hour every 10 minutes*”
  - Insert items into DB table as they arrive
  - Execute query every 10 minutes
- One query: slow and inefficient
  - Doesn’t reuse previous work
  - Can describe computation incrementally with recursion
- Computations
  - Can be complex (e.g., train LLM model with 1,000 GPUs)
  - Not always easily rewritten recursively
- Typically thousands of continuous queries



$$m_n = m_{n-1} + \frac{a_n - m_{n-1}}{n}$$

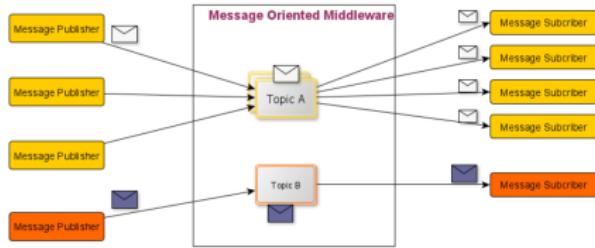
# Pub-Sub Systems

- In modern architecture, complex distributed systems are built with **small and independent building blocks**
  - E.g., serverless, micro-services
  - E.g., Uber
  - Facilitate decoupling and allow systems to evolve independently
  - Allow scalability and flexibility
- **Publish-subscribe systems**
  - Aka pub-sub, message queues, message brokers
  - Provide communication and coordination between building blocks
  - Focus on data dissemination
    - Typically no complex queries
    - Topics categorize messages and events
  - E.g., AWS SQS, AWS Kinesis, Apache Kafka, RabbitMQ, Redis, Celery, JBoss



# Pub-Sub Systems

- **Publishers**
  - Send messages or events
- **Subscribers**
  - Consume messages
- **Message broker**
  - Manage message flow between publishers and subscribers
- **Design parameters**
  - Events distributed
    - Topics, by subscription
  - Event distribution method
    - Push, pull
  - Subscriber interest in topics
  - Delivery semantics
    - At-most once
    - At-least once
    - Exactly once



# Delivery Semantics

- **At-most once**

- Message may be lost, not redelivered
- Pros
  - High-performance
  - Small implementation overhead
  - Easy to implement: “fire-and-forget”
- E.g., monitoring metrics allow small data loss



- **At-least once**

- Deliver message more than once, no message lost
- Handle transport message loss
  - Keep state at sender
  - Acknowledge state at receiver
- Works if
  - Data duplication is acceptable
  - Deduplication (e.g., storing key-value)
  - Idempotency



- **Exactly once**

- Every message sent once
  - Friendly for downstream consumers
- Difficult to implement
- Two Generals' Problem



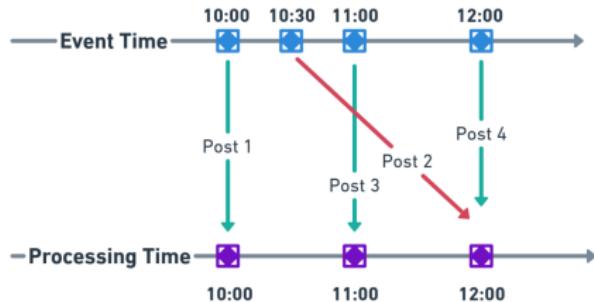
SCIENCE  
ACADEMY

Difficult to implement

• Two Generals' Problem

# Event vs Processing Time

- In both streaming and pub-sub architectures
- **Event time**
  - Time when each record is generated
- **Processing time**
  - Time when each record is received
  - Ingestion vs processing time: when events are received vs processed
- **Problems with events**
  - Out of order
  - Tardiness
  - How long to wait for late data?
    - In an asynchronous system, never sure all data has arrived
    - Use bounds on delay
    - Assume data arrives every minute
  - Recompute once late data arrives? If not, drop late data?



# Apache Streaming Zoo

---

- Many different streaming frameworks
  - Apache Apex, Apache Beam, Apache Flink, Apache Kafka, Apache Spark, Apache Storm, Apache NiFi
- Use cases
  - Real-time analytics
  - Online machine learning
  - Continuous computation
  - ETL processes
  - Data pipeline processing
  - Messaging
  - Log aggregation
- Different solutions to the same problem
  - Batch vs streaming
  - Delivery semantic type
  - Computing vs messaging/pub-sub
  - Throughput vs fault-tolerance
  - Supported languages
- Built simultaneously at different companies, then open-sourced

# Apache Storm

---

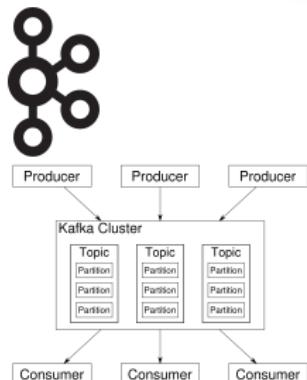
- Open-source distributed real-time computation system
  - Acquired and open-sourced by Twitter
- Horizontal scalability: add machines to handle increasing data
- Fault tolerance: **at-least-once** processing, automatic task restarts, workload redistribution
- Directed acyclic graph (DAG) with:
  - spouts (data sources)
  - bolts (processing units) as vertices
  - data streams as edges
- Suitable for complex data processing workflows with multiple stages and parallelism



APACHE  
STORM™

# Apache Kafka

- Open-source distributed streaming platform
  - Developed at LinkedIn, open-sourced in 2011
- Producers, brokers, consumers, topics, partitions
- Persistent storage, data replication across brokers
- High-throughput, low-latency messaging
- Fault tolerance: broker replication, automatic recovery
- Message delivery semantics:
  - **At-least-once, at-most-once, exactly-once**
- Kafka Connect: Integrate with data sources and sinks
- Kafka Streams: Stream processing library on Kafka



# Apache Flink

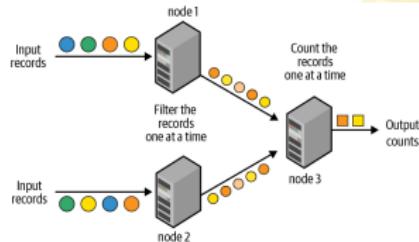
---

- Open-source, distributed data processing framework
- Focus on stateful computations over data streams
- Scalability: Horizontal scaling across large clusters
- Fault tolerance:
  - **Exactly-once** processing semantics, checkpointing, state management
- Batch processing support: **unified API for stream and batch processing**
- Flexible windowing: time-based, count-based, session windows
- Deployment options: standalone, YARN, Mesos, Kubernetes, cloud environments



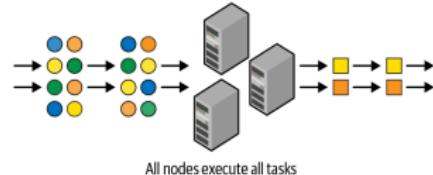
# Record-at-a-time Processing

- Implemented in Apache Kafka
- Goal: handle endless data stream
- **Multiple-node distributed processing engine**
  - Map computation on DAG of nodes
  - Each node continuously
    - Receives one record
    - Processes it
    - Forwards to next node
- **Pros**
  - Very low latencies
    - E.g., less than msec
- **Cons**
  - Inefficient node failure recovery
    - E.g., need failover resources/redundancy
  - Straggler nodes (slower than others)



# Micro-Batch Stream Processing

- Aka DStreams
- Implemented in Spark Streaming
- Computation as a continuous series of batch jobs on small chunks of stream data



- E.g., 1 second
- Process each batch in the Spark cluster in a distributed manner

## • Pros

- Recover from failures and stragglers using task scheduling
  - E.g., schedule same task multiple times
- Deterministic tasks
  - Exactly-once processing guarantees
  - Consistent API: same functional semantics as RDDs
  - Fault-tolerance

## • Cons

- Higher latency
  - E.g., seconds

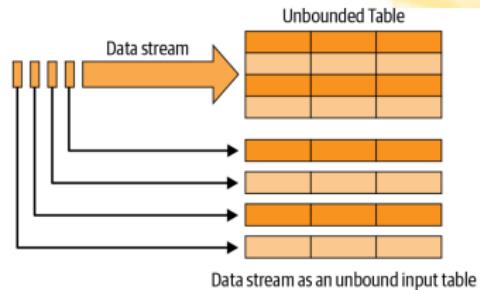
# Spark Micro-Batch Processing: Cons

---

- Line between real-time and batch processing blurred
  - Application computing data hourly: stream or batch?
- No single API for batch and stream processing
  - Same abstractions (RDD) and operations
  - Rewrite code with different classes
- No support for event-time windows
  - Operations defined by processing time
  - No support for tardy data
- Spark replaced DStreams with Structured Streaming in v3
  - Supports micro-batch and continuous streaming
  - Closer batch vs streaming API

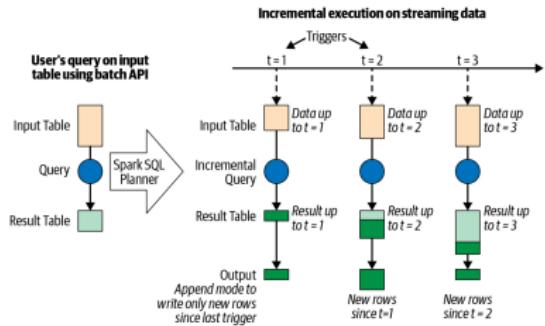
# Spark Structured Streaming

- New approach used by Spark
- Goal: write stream processing as easy as writing batch pipelines
  - Single unified programming model
  - Use SQL or DataFrames on stream
- Handle automatically
  - Fault tolerance
  - Optimizations
  - Incremental computation
  - Tardy data
- **Data abstraction**
  - Batch applications: table (DataFrame) is the abstraction
  - Structured Streaming: table is an unbounded, continuously appended table
    - New record becomes a new row appended
    - At time  $T$ , it's like a static dataframe with data until  $T$



# Incrementalization

- Automatically detect state to maintain
  - Build DAG of computation
  - Express output of graph at time  $T$  in terms of graph at time  $T-1$
  - Cache results
- Developers specify triggers to update results
- Incrementally update result with each record arrival



# Triggering Modes

---

- Indicate when to process newly available streaming data
  - **Default**
    - Process micro-batch after previous completes
  - **Trigger interval**
    - Specify fixed interval for each micro-batch
    - E.g., "every 10 minutes"
  - **Once**
    - Wait for external trigger
    - E.g., "at end of day"
  - **Continuous (experimental)**
    - Process data continuously
    - Not all operations available
    - Lower latency

# Saving Data

---

- Each time result table updates, write to external file system (e.g., HDFS, AWS S3) or DB (e.g., MySQL, Cassandra)
  - **Append mode**
    - Append new rows since last trigger
    - Use when existing rows don't change
  - **Update mode**
    - Write updated rows since last trigger
    - Update in place
  - **Complete mode**
    - Write entire updated result table
    - General but expensive

# Spark Streaming “Hello world”

- *lines* looks like an RDD but it's a *DataStreamReader*
  - Unbounded DataFrame
  - Set up reading but doesn't start reading
- *words* split data in words
- *counts* is a streaming DataFrame
  - Running word count
- Stateless transformations don't require maintaining state
  - E.g., *select()*, *filter()*
- Stateful transformations
  - E.g., *count()*
- Define how to write processed output
  - Where to write (e.g., *console*)
  - How to write (e.g., *complete* for updated word counts)
- When to trigger computation
  - E.g., every 1 second
- Where to save metadata for:

```
from pyspark.sql.functions import *
spark = SparkSession...
lines = (spark
    .readStream.format("socket")
    .option("host", "localhost")
    .option("port", 9999)
    .load())

words = lines.select(split(col("value"), "\\s").alias("word"))
counts = words.groupBy("word").count()
checkpointDir = "..."
streamingQuery = (counts
    .writeStream
    .format("console")
    .outputMode("complete")
    .trigger(processingTime="1 second")
    .option("checkpointLocation", checkpointDir)
    .start())
streamingQuery.awaitTermination()
```

